

Sunlumo AI Sales Challenge: Proposal

The opportunity

The defining failure mode of autonomous B2B sales in 2025 is well-documented: generic AI copy, hallucinated personalisation, burned deliverability, and damaged brand reputation. The structural cause is the same across every major platform; campaigns deploy first and learn later. For US SaaS teams optimising for volume that is tolerable. For a European battery distributor selling into a fragmented installer market where partner trust matters more than reply count, it is the wrong architecture entirely.

Our proposal is built on a different principle. **Before any message reaches a real installer, it is tested against synthetic installer personas grounded in your own historical data.** No major sales platform on the market today (Artisan, 11x, Clay, Apollo, Outreach, Salesloft, Cognism, Amplemarket, Lemlist, Humanlinker) ships pre-deployment campaign simulation. That is the gap we build into.

What we are building

A simulation-first agentic sales operating system, purpose-built for the European distributor-to-installer motion. It runs as a small set of specialised AI agents (research, simulation, outreach, feedback) under a central supervisor, exposed through a visual workflow canvas where your team designs and approves campaigns without writing code. Every campaign passes through a mandatory simulation step that forecasts reply rates, channel performance, and conversion before deployment. Real outcomes flow back through Bayesian updating to tighten the model, campaign by campaign.

Compliance presets (UWG, GDPR), native multilingual generation, a solar-vertical knowledge layer covering product compatibility, subsidy programmes, and certification requirements, plus channel-aware routing across email, WhatsApp Business, and human-in-loop LinkedIn are built in. US-built generalists are structurally disqualified by language, compliance, and vertical depth; European generalists lack the vertical layer.

Methodology in four steps

1. **Behavioural prior extraction.** Ingest your CRM, email history, support logs, and webshop data. Extract empirical base rates per installer archetype: reply rates, objection patterns, channel preferences, deal cycles. These are hard constraints, not generated by the model.

2. **Constrained persona construction.** Each archetype combines a factual demographic shell (from trade registers and enrichment sources), an empirical behavioural cage (from your data), and an LLM-generated conversational layer that operates only inside those constraints. The model decides how a persona phrases an objection; it does not decide whether they object.
3. **Simulation as aggregation.** Multi-step sequences are run against the persona cohort, and outcomes are counted, not invented. A predicted reply rate is the descriptive average across simulated runs, with a confidence interval, never a number the LLM has guessed.
4. **Retrodiction and live updating.** Before any forward prediction is trusted, the system reconstructs your historical campaigns and is measured against actual outcomes. After deployment, every real result tightens priors via Bayesian updating, with calibration tracked on a dashboard you can audit.

Sequencing: warmest signals first

The pilot starts deliberately with the most legally and commercially safe pocket: existing-customer reactivation triggered by intent signals (webshop, deanonymisation). Trade-show follow-up and subsidy-event campaigns come next; cold outbound to net-new prospects is sequenced last and only after warm-pipeline economics are proven. This ordering is a design choice; most competing pitches will invert it and put your brand at risk in the first quarter.

What we will not do

- No mass cold outbound in regulated markets before consent infrastructure is in place.
- No LinkedIn automation that violates platform terms; we use Sales Navigator manually for high-value accounts.
- No simulation predictions claimed as accurate before historical retrodiction has passed our internal acceptance threshold.
- No "AI replaces your team" framing; the system augments your KAM team and escalates ambiguous replies with full context.

Indicative LLM and infrastructure costs

The system uses frontier reasoning models for simulation and outreach copy, smaller models for batch enrichment, and self-hosted embeddings for retrieval. Estimated pass-through costs, billed at provider price plus a small operations margin:

Stage	Monthly LLM cost (indicative)	Supporting infra
Phase 0 MVP, single-market pilot	€500 to €1,500	€400 to €800
Full DACH production	€2,000 to €5,000	€800 to €1,500
Pan-European, multiple active campaigns	€5,000 to €10,000	€1,500 to €2,500

These are planning ranges, not commitments. A live cost dashboard is shared from week 1, and we will tune simulation depth and model mix to keep spend predictable as volume scales.

Timeline: first month

The engagement opens with a 30-day proof-of-life. A 14-day milestone sits inside it as the first hard checkpoint.

Week	Milestone	What you see
1	Data connectors live (CRM, email, webshop, deanonymisation).	Last 12 months of data ingested into the prior-extraction pipeline.
2 (14-day checkpoint)	First archetypes built; simulation engine v0.1; first historical retrodiction.	Live screen-share showing predicted-vs-actual gap on past campaigns.
3	Visual workflow canvas with one production template.	Your team designs and simulates a campaign without engineering help.
4	First production deployment to a small consented cohort.	Live dashboard: predicted reply rate alongside real-time actuals.

By the end of month 1 you have working infrastructure ingesting your data, a calibrated simulation engine, a canvas your team can operate, and a first real campaign deployed with measurement live.

Commercial structure

The contract is phased and milestone-gated, sized inside the range named in the brief.

- **Phase 0, build and calibrate (months 1 to 3):** fixed scope. Covers infrastructure, prior extraction, simulation engine, canvas, first templates, retrodiction validation, and a controlled A/B against your current best-practice sequence.
- **Phase 1, DACH scale (months 4 to 6):** regional rollout, trade-show and subsidy-trigger workflows, public-facing knowledge-graph asset.
- **Phase 2, expansion (months 7 to 12):** additional European markets, each onboarded as a separate calibration run with its own KPI commitment.
- **Phase 3, long-term partnership (month 13 onward):** renewable annual engagement covering model maintenance, new channels, and continued expansion.

Performance is tracked against qualified partner registrations and resulting orders, measured against your current cost baseline. The simulation layer is designed to outperform untested cold outreach, but the magnitude is something we measure together rather than promise upfront. Phase gates between each stage allow either side to review the engagement openly if a phase misses its KPI.

What we are asking for

A short decision window and a Phase 0 contract. We start the day Sunlumo signs. By the 14-day checkpoint you will have working infrastructure, a calibrated simulation engine, and a real campaign in flight; every claim above becomes empirically testable inside the first month.

The brief asked for builders, not retrofitters. That is what we are: two AI systems engineers, one shared channel with your team, weekly demo cadence, no agency layer, no offshore handoff.

This is not "an AI SDR like 11x, in German." It is the first simulation-tested, vertically specialised B2B sales system for the European distributor-to-installer layer, with Sunlumo as the reference case.